

# Predicting phenotypic traits of microbes from their genome

Fred Farrell<sup>1</sup>, Orkun S. Soyer<sup>1\*</sup>, Christopher Quince<sup>2\*</sup>,

**1** School of Life Sciences, University of Warwick, Coventry, United Kingdom

**2** Warwick Medical School, University of Warwick, Coventry, United Kingdom

\* Joint corresponding authors: c.quince@warwick.ac.uk, o.soyer@warwick.ac.uk

## Abstract

Predicting function from genomic information remains an unresolved problem. One that is particularly relevant in the context of microbial communities research, where we are increasingly able to use metagenomics to obtain novel genomes from a variety of natural and man-made environments such as the ocean, human gut, and biotechnological reactors. Here, we develop and test different machine learning based classifier models to assign functions to metagenome assembled genomes (MAGs). We have trained and cross-validated these models using a database of 84 phenotypic traits linked to 9,407 available, fully-sequenced genomes. We found that different classifier models perform differently and when using different feature vectors. The best model performance resulted when we used a lasso logistic regression model with KEGG ortholog usage frequency as the decision feature. This model was able to classify 65 of these traits with greater than 90% cross-validated AUROC score. We show the utility of this classifier model by assigning traits to novel MAGs from three different metagenomic studies focusing on different environments.

## Introduction

Predicting phenotype from genotype remains one of the major challenges in biology [1, 2]. Addressing this challenge is particularly relevant for understanding microbial communities, the study of which has been boosted by an increased ability to resolve sequence data directly from communities. Technological improvements in DNA sequencing have led to an explosion in the amount of such data generated. In the context of microbial ecology, large-scale metagenomic studies such as the Human Microbiome Project [3], the Earth Microbiome Project [4] and the Tara Oceans Project [5] have systematically sequenced the microbial communities in a huge variety of environments at great depth. Amplicon sequencing, such as of the 16S rRNA gene, allows detailed study of the taxonomic makeup of these communities, while shotgun metagenomic sequencing allows characterisation of all genes present in an environment. Increasing depth of coverage and improvements in genome binning algorithms for clustering contigs into genomes, in particular the use of differential coverage across different samples [6, 7], are allowing more and more full and partial genomes to be assembled from shotgun metagenomic studies. Many of these organisms are novel and uncultured, having never been studied in a lab. A recent metagenomic study on aquifer systems [8], for example, reconstructed 2540 separate high-quality, near-complete genomes, and claimed to have discovered an astonishing 47 new phylum-level lineages among them.

Converting this exponentially growing sequence data into functional understanding of microbial communities requires us to determine physiological functions from it [9]. This would allow the inference of key functions in microbial communities found in human and animal guts, soil, and the oceans, and how these functions change over ecological conditions and with time [9]. In turn, this ability, could allow us to discern ecological adaptations in environmental microbial communities, as well as to achieve functional mechanistic models of stability and function [10]. It would be a key step towards managing communities underpinning human and animal health and biogeochemical cycles.

Efforts to achieve phenotype-genotype mapping from environmental sequence data has so far mostly focussed on phylogenetic assignments using the 16S rRNA gene. This highly conserved gene can provide a phylogenetic assignment at the species (or higher) level, which can then be used to infer general functional traits. While this approach has been commonly used to study ecological distribution of microbial functions e.g. [11–13], its premise of a direct association of function with phylogenetic assignment (i.e. ‘functional coherence of microbial taxa’) is questionable (e.g. [14]). The level of taxonomic coherence of function is not clear even for strains of the same species, where functions can show high variability either due to a few genetic changes or even regulatory changes [2]; as seen for example among *Escherichia coli* strains [15]. It is also a common problem that when certain taxonomic groups in a microbial community are found to show direct associations with certain ecological factors (or health state of an host), these groups are phylogenetically so broad that assigning specific functions to them is hard or impossible [16]. Indeed, a study of specific functional traits across microbial taxa has found that many of these traits are dispersed across the phylogenetic tree [2,17]. Even where a specific functional trait is taxonomically coherent, the phylogenetic approach is limited by our ability to assign taxonomy based on the 16S rRNA gene. The extensive accumulation of metagenomics data indicates that we have sampled only a fraction of microbial diversity, and it is not uncommon for such data to result in many unassigned taxa, as for example by the aquifer study described above which described dozens of apparently novel phyla.

The scope to perform functional characterisation of environmental samples has expanded with the advent of metagenomics sequencing. This technology has allowed development of several bioinformatics pipelines that can go from raw metagenomic sequence data to predicted (i.e. binned) genomes. These metagenome assembled genomes (MAGs) are then analysed for the presence of specific functional traits (e.g. [18]). The functional annotation steps in existing bioinformatics pipelines usually considers presence or absence of specific genes in a MAG, and associated with known metabolic pathways, as identified for example in databases such as the Kyoto Encyclopedia of Genes and Genomes 22 (KEGG) [19], the Pfam database of protein families [20] and the NCBI COG database of orthologous genes [21]. This approach circumvents the problems of taxonomic coherence and 16S rRNA gene based assignment to taxa, yet relies heavily on existing categorization of genes into functions as done in the above databases. While these functional gene groupings are mostly based on accumulated knowledge and experimental data on metabolic pathways, they might miss the full set of genes associated with a given function and do not consider functions that cannot be assigned to a few genes or seemingly well-organised pathways. One route to overcome such limitations is to develop extensive databases of phenotypic traits of microbes without necessarily using a pathway-centric view. These functional assignments could then serve as a source to apply statistical approaches to ‘learn’ genetic drivers of those functions using genomes of associated microbes. Efforts in this direction have recently resulted in the compilation of literature-based assignment of functions in microbes, either covering a large selection of functions and

organisms [22,23] or specific ones such as methanogenesis [24]. The FAPROTAX database [22], which we focus on here, is based on an extensive survey of the scientific literature. The aim of creating this database was to allow microbes found from 16S rRNA amplicon sequencing to be assigned into functional and metabolic groups, so that functional variation across environments could be studied and compared to taxonomic variation. The authors found that the abundance of functional groups was strongly influenced by environmental conditions in a variety of ocean environments [12]. The bulk of the classifications in the FAPROTAX database come from *Bergey's Manual of Systematic Bacteriology* [25], and it currently contains 84 phenotypic traits associated with 4600 microbial taxonomic groups.

Here, we use this accumulating data on phenotypic traits to develop an algorithmic approach to create functional assignments of MAGs. To do this, we combined the FAPROTAX functional trait database with 9407 genomes downloaded from the NCBI, and used machine learning based approaches to train statistical models able to infer an organism's phenotypic traits from their genome. For model construction and training we compared the use of gene orthologs from the KEGG database and Pfam domains. We show that the resulting classifiers perform better than simple taxonomic assignments of trait, and reveal both known and new genetic drivers of specific traits. Using the resulting classifiers for over 80 traits, we were able to analyse three recent, large-scale metagenomics datasets from three diverse environments: anaerobic digesters, ground water aquifers and the ocean. The classifier-based functional analysis of these datasets revealed significant differences in functional properties between environments and conditions. This work was inspired by a recent software framework, Traitar [26], which uses SVMs to predict microbial traits based on genomic information in the form of copy numbers of Pfam families. Our work differs from Traitar in the use of the highly detailed FAPROTAX database. Traitar utilised the Global Infectious Disease and Epidemiology Online Network (GIDEON) [27] for its phenotypic annotations, and was therefore biased toward pathogenic traits; we instead focus on traits associated with metabolism and environmental niche.

## Materials and methods

### Databases and preparation of training data

To train the models, we utilized the combination of the recently published FAPROTAX database of microbial phenotypes and the NCBI genome database. We downloaded all prokaryotic genomes classified as 'full' from the NCBI Genome database. We used the taxonomic information available from NCBI to assign them phenotypes using the script 'collapse\_table.py' which comes as part of the FAPROTAX database [22]. We then called genes in these genomes using Prodigal [28]. We annotated the resulting inferred coding DNA sequences (CDS) both by aligning against the KEGG database using Diamond BLASTP [29] and by searching with hmmer3 [30] against Pfam [20] all with standard settings. The result is a matrix of organisms and their copy numbers of either KEGG orthologs or Pfam domains. We have found that using gene copy number rather than simple presence/absence significantly increases classifier performance. The scripts we used to download and process the genomes are available at <https://github.com/chrisquince/GenomeAnalysis.git>.

### Statistical modelling

To model the link between genotype, in the form of KEGG ortholog copy numbers or Pfam protein families, and phenotypes (i.e. class targets) as represented in the

FAPROTAX database, we used a variety of machine learning techniques. These are all different ways of learning the relationship between the features (gene copy numbers) and targets (biological traits) from the training data of 9407 NCBI genomes from unique species. To train the algorithms, we first split this data into a training set (75% of the genomes) and a test set (25%), by random sampling. The algorithms were then trained on the training data, and their performance tested on the unseen test data genomes, to check that the relationships learned are generalisable (i.e. that the algorithms have not ‘overfit’ the training data). Below, we describe the algorithms used.

### Logistic regression

We found logistic regression, a commonly used linear model for classification problems, [31,32] to be an effective strategy in our case. We scaled all input features to have mean zero and variance one before performing the regression. Since the number of KEGG orthologs (features) was somewhat larger than the number of training examples, overfitting, whereby the model classifies on features of the training set which are very specific to it, was potentially a problem. To alleviate this, we used logistic regression with an  $\ell_1$  penalty term, also known as LASSO logistic regression [33], whereby parameters are penalized in such a way that only a few of the features have a nonzero weight. In detail, the method involves adding a penalty term equal to the  $\ell_1$ -norm of all of the coefficients of the regressor, thereby penalising nonzero terms, to create an optimisation problem defined by;

$$\min_w \left[ \|w\|_1 + C \sum_{i=1}^n \log (\exp(-y_i(X_i^T w)) + 1) \right] \quad (1)$$

where  $w$  is the vector of regression weights,  $X_i$  are the feature vectors of each sample,  $y_i$  the classification targets,  $n$  is the number of data points, and  $C$  is a parameter defining the (inverse) strength of the regularization. This method of regularization is often useful in cases where the number of features is large (similar to or larger than the number of training examples), as most of the features are not used in the classification task.

### Random forests

We also used the random forest algorithm, a popular machine learning method which can be applied to both regression and classification problems, which is simple to use, fast and performs fairly well on a wide variety of problems [34]. The random forest is an ‘ensemble’ method, using a collection of slightly randomized classifiers, the results of which are averaged to produce a prediction. This helps to avoid overfitting. A random forest is an ensemble of so-called decision trees. A decision tree is a model which learns to split up training examples into sets according to their feature values, with the aim of separating the target classes. They have the advantage of being invariant under scaling of features and adding of irrelevant features, this last feature being useful in our case where the number of features is very large and many are irrelevant to the classification task; they can also learn more complex relationships between variables than a linear model such as logistic regression. However, an individual tree tends to overfit the training data. A random forest trains a large number of such trees on random subsets of the features and combines these predictions by averaging, avoiding overfitting and much improving performance.

### Support vector machines

Finally, we used support vector machines (SVMs) [31]. An SVM essentially tries to find surfaces in the high-dimensional feature space which separate the different classes as

well as possible, and with as wide a margin as possible between the surface and the examples. These surfaces can be either linear or non-linear (if a non-linear kernel is used); they are therefore capable of capturing complex non-linear relationships between features and targets. They can also include regularization terms as in logistic regression, to reduce overfitting.

164  
165  
166  
167  
168

## Metrics and classifier performance

169  
170  
171  
172  
173  
174  
175  
176  
177  
178  
179  
180

Since many of the classes which we are attempting to predict are highly unbalanced (e.g. of the 9407 unique species with full genomes in the NCBI database only 83 are hydrogentrophic methanogens), simple classification accuracy is not a very useful measure of classifier performance. Predicting all labels as negative in the above example would give an accuracy of 99.1% despite not being a useful classifier. We therefore need a metric which can take into account class imbalance. We use the area under the ROC (Receiver Operating Characteristic) curve, which is a graph of true positive rate against false positive rate as one varies the cutoff in probability for making a positive prediction [35, 36]. An AUROC score much greater than 0.5 (the score for random predictions) indicates a good classifier. In particular, a score of 1 indicates that all positive cases have been assigned a higher probability than all negative cases.

## Prediction of MAG phenotypes

181  
182  
183  
184  
185  
186  
187  
188

Once classifiers have been trained on the NCBI data, it is possible to use them to make predictions about unseen genomes, such as MAGs generated from shotgun sequencing studies. The MAGs must first be processed to give a matrix of the KEGG ortholog copy numbers associated with them, using the same pipeline as applied above to the NCBI genomes. These matrices are then used as input into the classifiers to produce a matrix of MAGs and their predicted traits, which can be either presence/absence predictions or probabilities.

## MAG collections

189  
190  
191

We applied our classifier to three separate collections of MAGs from three different studies:

- Tara Oceans MAG collection: This comprised a subset of 660 MAGs from the collection of 957 non-redundant MAGs generated from the Tara Oceans microbiome in Delmont *et al.* [37]. These 660 MAGs were those which had at least 75% of the 36 single-copy prokaryotic core genes identified in Alneberg *et al.* [6] in a single-copy and can thus be considered reasonably complete and pure prokaryotic genomes. The Tara Oceans microbiome survey generated 7.2 terabases of metagenomic data from 243 samples across 68 locations from epipelagic and mesopelagic waters around the globe [38], Delmont *et al.* extracted their MAGs from a subset of 93 of these samples, 61 surface samples and 32 from the deep chlorophyll maximum layer. Therefore these MAGs represent a substantial sample of planktonic microbial life.
- Anaerobic digester (AD) MAG collection: This comprised a collection of 153 MAGs that were constructed by co-assembly and binning of 95 metagenome samples taken from three replicate laboratory anaerobic digestion (AD) bioreactors converting distillery waste into biogas. They were assembled with Ray using a kmer size of 41 and all 186,081 contig fragments greater than 2 kbp in length were clustered by CONCOCT [6] generating a total of 355 bins of which 153

had 75% of single-copy core genes present in a single copy. These were considered high quality enough to be used in this analysis.

- Candidate phyla radiation (CPR) MAG collection: This collection of 581 MAGs is a subset of 797 MAGs provided by the authors of Brown *et al.* 2015 [39]. They comprise members of the Candidate phyla radiation (CPR) assembled from ground water enriched with acetate.

## Results

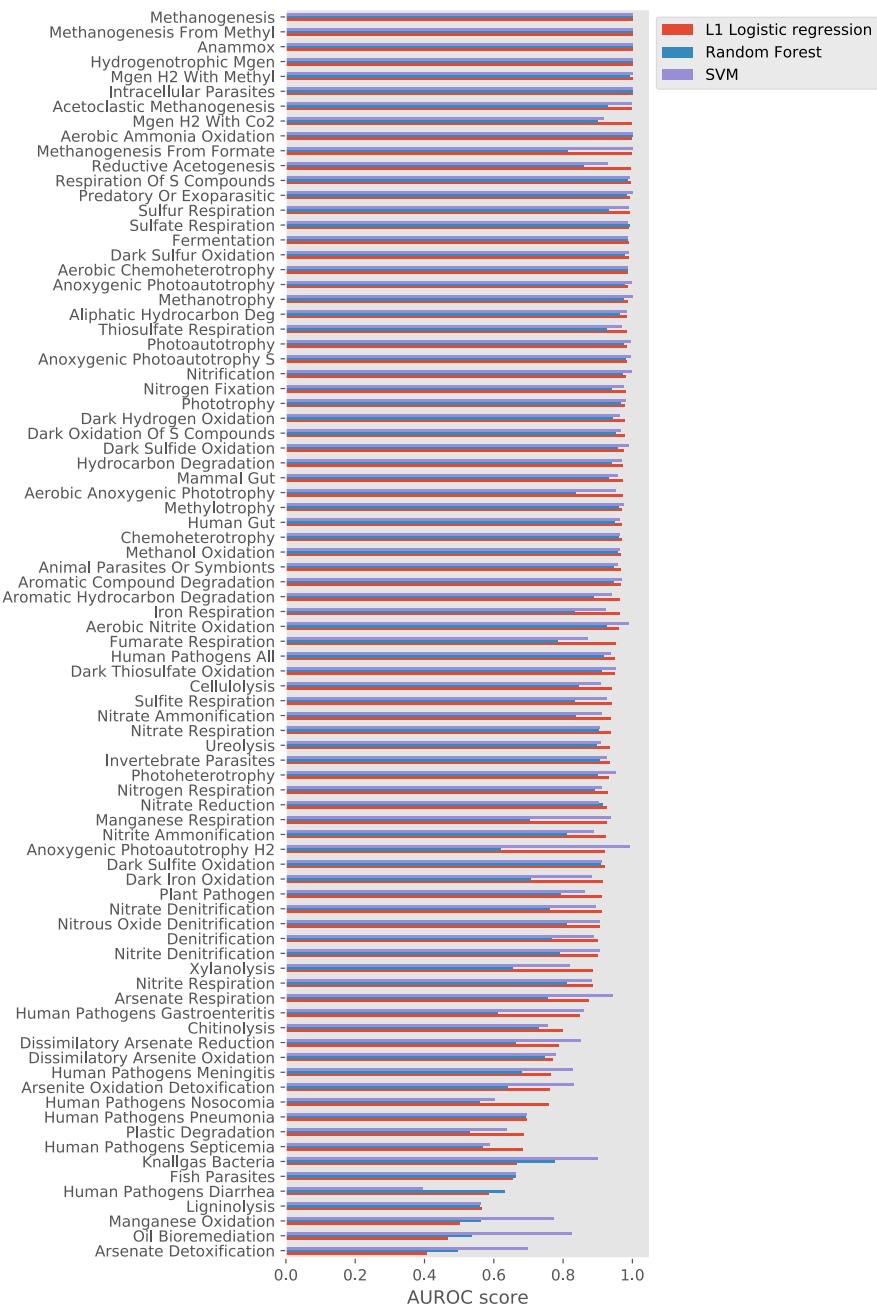
### Classifiers based on genomic features achieve over 90% accuracy in assigning phenotypic traits

To create a phenotypic trait assignment for MAGs, we have created different machine learning classifiers and that take as input genomic features. Specifically, we build classifiers based on genes and their copy numbers (see Methods). Genes were identified from the genome using existing bioinformatics pipelines and the databases of KEGG orthologs (KOs) [19] or Pfam gene families [20]. Using the resulting gene and copy number information as ‘features’ of a genome, we then trained classifiers on these features to assign specific phenotypic traits to those genomes. We used as a training platform the known phenotypic traits, as compiled in the FAPROTAX database [22], and over 9000 fully annotated genomes, as listed on the NCBI database. We used this general approach to develop different machine learning based classifiers (see Methods).

To evaluate the performance of the different classifiers, we used the Area Under Receiver Operating Characteristic (AUROC), a common measure to evaluate classifiers (see Methods). The AUROC scores were calculated using k-fold cross-validation with  $k = 5$ , i.e. the data was split into training and testing sets 5 times, in such a way that each training example was in the test set once, and the prediction for each data point when it was in the test set was used. The results are shown in Figure 1 for the three classification algorithms,  $\ell_1$ -regularized logistic regression, the random forest and a linear SVM. The regularized logistic regression outperforms the random forest in assigning most, though not all, of the phenotypic traits assessed. The average score over all traits for logistic regression is 90.1% (versus 84.5% for the random forest), and 65 traits have a score greater than 90%, with 45 higher than 95%. The performance of the SVM and the logistic regression are similar, although they do differ significantly for some traits. The mean score of the SVM is slightly better, at 90.8% vs. 90.1%. This difference is not statistically significant (paired t test,  $p=0.71$ ), and since logistic regression is easier to interpret and much more computationally efficient, we decided to focus on it for the rest of the paper.

Additionally, we can compare the results we obtain using the two different ways of obtaining genes (and gene copy numbers) from genomes, namely KOs and Pfam families, see Figure S1. The results using the two schemes are rather similar, though there are some traits where one outperforms the other; this may reflect better coverage of the genes involved in the trait in a particular scheme. On average, KO performs better, with a mean score of 90.1% versus 84.9% for Pfam ( $p < 0.001$ ), and we therefore concentrate on the approach of obtaining genes from genomes using KOs for the rest of the analysis presented below.

It is instructive to compare the performance of the classifiers to the use of KEGG modules, where an equivalent module exists for that trait, i.e. compare the performance to a ‘classifier’ where an organism is judged capable of a trait if it has a complete KEGG module for that trait. Table 1 shows the results of this comparison for three FAPROTAX traits with corresponding KEGG modules. Note that the KEGG module



**Fig 1. Overall performance of classification algorithms.** The AUROC score on each classification task (each trait) is shown for three classification algorithms:  $\ell_1$ -regularized logistic regression, the random forest and a linear support vector machine (SVM). Traits are ordered by the logistic regression score.

method does not require training, so the metrics are over the entire NCBI dataset, whereas for the classifier they are only for the held-out test set. Also, the former method gives only presence/absence of a trait rather than a probability, so the AUROC score cannot be calculated, so we use alternative metrics based on classification: the  $F_1$  score and the confusion matrix.

257  
258  
259  
260  
261

| module                          | KEGG modules |  | Classifier |   |
|---------------------------------|--------------|--|------------|---|
|                                 | F1           | confusion matrix   | F1         | confusion matrix                                    |
| sulfate respiration             | 0.84         | $\begin{pmatrix} 9197 & 53 \\ 6 & 151 \end{pmatrix}$     | 0.99       | $\begin{pmatrix} 2313 & 0 \\ 1 & 38 \end{pmatrix}$  |
| nitrate respiration             | 0.14         | $\begin{pmatrix} 6821 & 2162 \\ 224 & 200 \end{pmatrix}$ | 0.622      | $\begin{pmatrix} 2237 & 9 \\ 54 & 52 \end{pmatrix}$ |
| hydrogenotrophic methanogenesis | 0.756        | $\begin{pmatrix} 9281 & 43 \\ 6 & 77 \end{pmatrix}$      | 0.923      | $\begin{pmatrix} 2331 & 0 \\ 3 & 18 \end{pmatrix}$  |

**Table 1. Comparison of classifiers to KEGG modules.** Table showing the performance of using KEGG module presence/absence against logistic regression classifiers for some traits where equivalent KEGG modules exist. Since the KEGG module approach does not give a probability, the AUROC score cannot be used, so the F1 score and confusion matrices are compared.

It can be seen that the logistic regression classifier does significantly better than KEGG modules in assigning these traits as they appear in the FAPROTAX database. This suggests that having the enzymes or proteins described in the KEGG module for a trait is not in fact a necessary or sufficient condition for actually performing that trait, and that other genes are more predictive. However, it is possible that the discrepancy is due instead to inaccuracy in the FAPROTAX database, e.g. species which do perform the traits being missed from the database and therefore getting flagged as false positives with the KEGG method. More work would be needed to fully exclude this possibility.

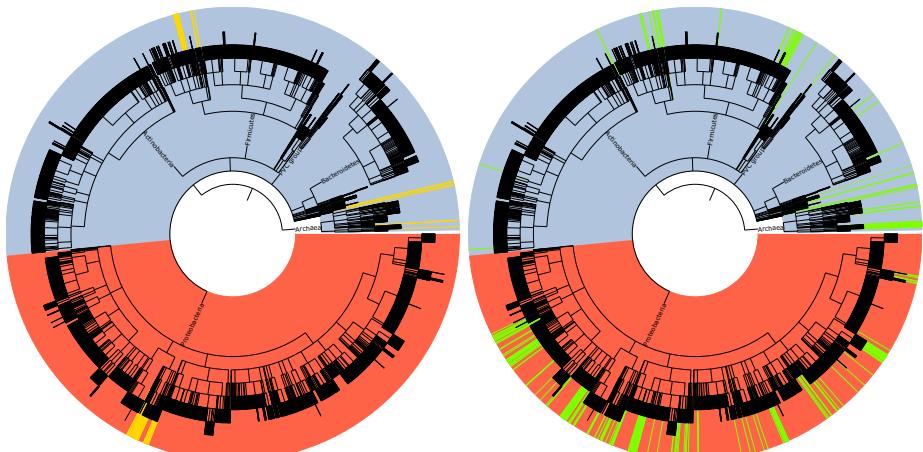
### Classifier performance is not due to taxonomic coherence

It is possible that the classifiers are simply picking genes that are irrelevant to a specific trait, but rather are present in taxonomically close species that all happen to perform a given trait. To explore this possibility, we first looked at how traits spread over the taxonomy. We found that this varies between traits (Figure 2), but as expected, organisms in the same taxa usually present similar traits.

To investigate this phenomenon and attempt to find orthologs with real causal associations with traits, we tried training a model on one taxa and testing its performance on the others. If a classifier can predict phenotype based on genes in a distantly-related, unseen set of organisms, it is likely the genes it is using have a real association with the trait. In particular, we tried training the logistic regression models on the Proteobacteria, a large phylum of bacteria, and testing on the rest of the taxonomy. Some traits did not have significant numbers of species in each of these sets; we used only traits with at least 5 species in the training set and 5 in the test, leaving 59 traits out of 84.

As might be expected, this ‘enforced’ ignoring of taxonomic information reduced the performance of the classifiers, compared to being trained on a random selection of species from throughout the prokaryotic part of the tree of life, see Figure 4b. However, for a significant number of traits the performance of the classifier is still fairly good, indicating an ability to make predictions which are generalizable to significantly different unseen groups of organisms. 19 traits have an AUROC score greater than 80%, and 9 greater than 90%.

Figure 4b shows a scatter plot of classifier complexity against performance for the models trained on Proteobacteria only. Compared to the use of all taxa in training, Figure 4a, it is notable that the group of classifiers achieving high accuracy while using a lot of genes is gone: traits such as fermentation and nitrate reduction, which were in this group of classifiers, are now much less accurate. Classifiers which work well in the cross-taxon case all use a relatively small number of genes, less than 150 or so. This suggests that the classifiers using a large number of genes to make predictions in the randomized case may have been using a range of genes found in different closely-related



**Fig 2. Taxonomic distribution of metabolic traits.** Taxonomic trees of all prokaryotic NCBI species with full genomes. For training the cross-taxa verison of the classifier, only the Proteobacteria (red section of the tree) were used, and the models were tested on the rest of the tree. Species capable of a) sulfate respiration and b) nitrate respiration are highlighted on the trees.

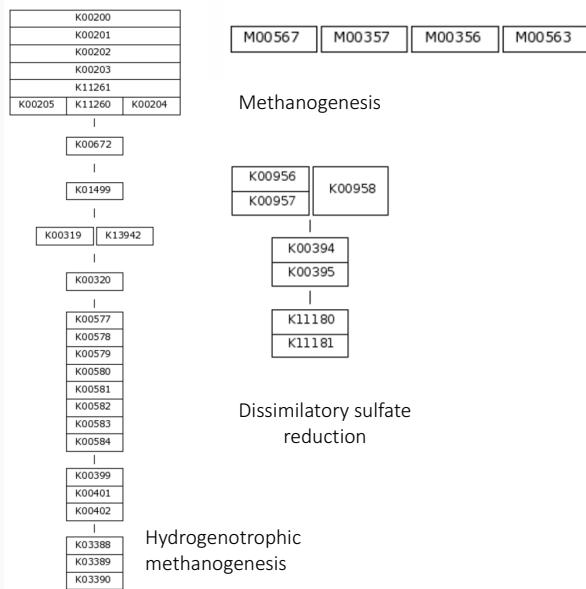
clusters of organisms which all have the target trait, but which may not have a causal relationship with the trait. 300  
301

### Resulting classifiers highlight the importance of specific and non-obvious gene combinations in defining trait assignments 302 303

The classifiers discussed above are built upon genomic features, in this case KO gene frequencies. Thus, we can analyse each classifier for each phenotypic trait and determine which genes underpin those traits. Table 2 gives the KOs that have non-zero coefficients in the resulting logistic regression classifiers for some example traits, for which the classifier had an AUROC score above 95%. Many of the KOs involved are those known to be involved in the associated trait, as one might have hoped. In particular, consider the prediction of methanogens, a relatively easy task since it is known that methanogens must possess the *mcrA* gene, this being a necessary and sufficient condition for methanogenesis [40]. Indeed, subunits of this gene have the highest weight, and a total of only 9 genes are used by the classifier. 304  
305  
306  
307  
308  
309  
310  
311  
312  
313

Looking at some more complex traits, for example sulfate respiration (i.e. disimilatory sulfate reduction to H<sub>2</sub>S), the model assigns a lot of weight to subunits of a quinone-modifying oxidoreductase, which is indeed associated with sulfur metabolism [41]. Interestingly, however, none of the genes picked out by the classifier are directly part of the metabolic pathway for this process as described in the KEGG module for disimilatory sulfate reduction, see Figure 3. The situation is similar with hydrogenotrophic methanogenesis, with classification mostly determined by components of energy-converting hydrogenases which are not directly part of the autotrophic methanogenesis pathway, along with *mcr* genes indicating that the microbe is a methanogen. 314  
315  
316  
317  
318  
319  
320  
321  
322  
323  
324  
325  
326  
327

Given this observation that different classifiers make use of different numbers of KOs, we further explored the relation between number of features (KO identity) used in a classifier and its performance (see Figure 4). This analysis demonstrates that there is indeed a correlation between these two variables, with some highly accurate classifiers built out of a large number of genes. However, there is also a noticeable cluster of traits 328



**Fig 3. KEGG modules for some traits.** Representations of the KEGG modules corresponding to the FAPROTAX traits shown in Table 2. Modules are organized into ‘blocks’ of orthologs, typically indicating a protein complex. Orthologs positioned next to each other are ‘options’, i.e. that section of the module is present if any of the adjacent blocks are present.

with high accuracy achieved with only a few genes (less than 100). These traits may be particularly interesting, as it is more likely that these small groups of orthologs are causally associated with the trait, rather than just being genes which typically occur in parts of the phylogenetic tree which have the trait and may or may not have any direct relation to it. This issue is explored further in the section on performance across taxa, below. Also, note that most of the traits that perform poorly, which typically use very few genes to classify, have very low support in the training data in terms of number of positive examples.

The ability of the classifiers to assign traits based on few and specific KOs shows that genomes with a specific trait must be clustering in the high-dimensional KO space. To explore this possibility, we have created an ordination plot (Figure S2) of all the species in the training and test sets using their KEGG ortholog copy numbers. That is, a dimensionality reduction algorithm, here stochastic neighbour embedding, has been applied to visualise variation in all KO copy numbers in two dimensions. Points are colored by whether they are true positive, true negative, false positive or false negatives under a particular classification task, here for sulfate respiration. It can be seen that the species performing this trait do tend to cluster together into a few groups in the KO space, allowing the algorithm to classify them mostly correctly.

### Trained classifiers can be applied successfully to analyse the phenotypic traits of MAGs

A major aim of training these classifiers is to explore the functional capabilities of novel genomes isolated from metagenomic studies. We therefore used the classifiers trained above to classify metagenomically assembled genomes (MAGs) from a few different environments. These were laboratory anaerobic digesters, the ocean (from the Tara

| Sulfate respiration |        |   |
|---------------------|--------|---|
| KO                  | Weight | Description   |
| K03421              | 0.263  | methyl-coenzyme M reductase subunit C   |
| K14109              | 0.226  | energy-converting hydrogenase A subunit R   |
| K14094              | 0.156  | energy-converting hydrogenase A subunit C   |
| K00401              | 0.131  | methyl-coenzyme M reductase beta subunit [EC:2.8.4.1]                                       |
| K14097              | 0.093  | energy-converting hydrogenase A subunit F   |
| K00440              | 0.091  | coenzyme F420 hydrogenase subunit alpha [EC:1.12.98.1]                                      |
| K06862              | 0.045  | energy-converting hydrogenase B subunit Q   |
| K16204              | 0.038  | seco-amyrin synthase [EC:5.4.99.52 5.4.99.54]   |
| K11099              | 0.033  | small nuclear ribonucleoprotein G   |
| K14098              | 0.027  | energy-converting hydrogenase A subunit G   |
| K09613              | 0.026  | COP9 signalosome complex subunit 5 [EC:3.4.-.-]   |
| K14093              | 0.022  | energy-converting hydrogenase A subunit B   |
| K08074              | 0.013  | ADP-dependent glucokinase [EC:2.7.1.147]  |
| K05181              | 0.013  | gamma-aminobutyric acid receptor subunit beta   |
| K09493              | 0.013  | T-complex protein 1 subunit alpha   |
| K06612              | 0.013  | alpha-N-acetyl-neuraminate alpha-2;8-sialyltransferase (sialyltransferase 8B) [EC:2.4.99.-] |
| K17278              | 0.011  | membrane-associated progesterone receptor component   |
| K02938              | 0.003  | large subunit ribosomal protein L8e   |
| K00442              | 0.003  | coenzyme F420 hydrogenase subunit delta   |
| K14096              | 0.003  | energy-converting hydrogenase A subunit E   |

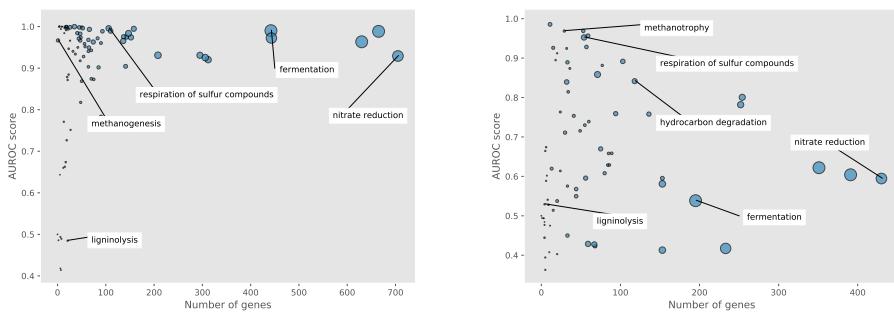
  

| Methanogenesis |        |   |
|----------------|--------|---|
| KO             | Weight | Description   |
| K03421         | 0.567  | methyl-coenzyme M reductase subunit C                               |
| K00400         | 0.276  | methyl coenzyme M reductase system; component A2                    |
| K00579         | 0.160  | tetrahydromethanopterin S-methyltransferase subunit C [EC:2.1.1.86] |
| K00399         | 0.081  | methyl-coenzyme M reductase alpha subunit [EC:2.8.4.1]              |
| K07463         | 0.023  | archaea-specific RecJ-like exonuclease                              |
| K17618         | 0.023  | ubiquitin-like domain-containing CTD phosphatase 1 [EC:3.1.3.16]    |
| K00401         | 0.022  | methyl-coenzyme M reductase beta subunit [EC:2.8.4.1]               |
| K09728         | 0.020  | uncharacterized protein   |
| K09613         | 0.002  | COP9 signalosome complex subunit 5 [EC:3.4.-.-]                     |

| Hydrogenotrophic methanogenesis |        |  |
|---------------------------------|--------|--|
| KO                              | Weight | Description  |
| K03421                          | 0.231  | methyl-coenzyme M reductase subunit C                  |
| K14109                          | 0.201  | energy-converting hydrogenase A subunit R              |
| K00401                          | 0.169  | methyl-coenzyme M reductase beta subunit [EC:2.8.4.1]  |
| K14098                          | 0.136  | energy-converting hydrogenase A subunit G              |
| K14097                          | 0.104  | energy-converting hydrogenase A subunit F              |
| K14093                          | 0.058  | energy-converting hydrogenase A subunit B              |
| K06862                          | 0.057  | energy-converting hydrogenase B subunit Q              |
| K14094                          | 0.049  | energy-converting hydrogenase A subunit C              |
| K17278                          | 0.043  | membrane-associated progesterone receptor component    |
| K08074                          | 0.042  | ADP-dependent glucokinase [EC:2.7.1.147]               |
| K00442                          | 0.032  | coenzyme F420 hydrogenase subunit delta                |
| K09613                          | 0.031  | COP9 signalosome complex subunit 5 [EC:3.4.-.-]        |
| K09493                          | 0.017  | T-complex protein 1 subunit alpha                      |
| K14099                          | 0.009  | energy-converting hydrogenase A subunit H              |
| K02938                          | 0.005  | large subunit ribosomal protein L8e                    |
| K00399                          | 0.001  | methyl-coenzyme M reductase alpha subunit [EC:2.8.4.1] |

**Table 2. Details of classifiers for specific traits.** Tables showing all the nonzero weights in the logistic regression models trained on three traits from the FAPROTAX database. Note that there are 9647 KEGG orthologs used in the models, so the vast majority of weights are set to zero in these models.



**Fig 4.** Scatterplots showing the AUROC score of the different classifiers plotted against the number of gene orthologs the classifier uses to make its predictions. Point size is proportional to the number of positive examples in the training set. Left: in the standard case. Right: in the cross-taxa case with training on Proteobacteria only.

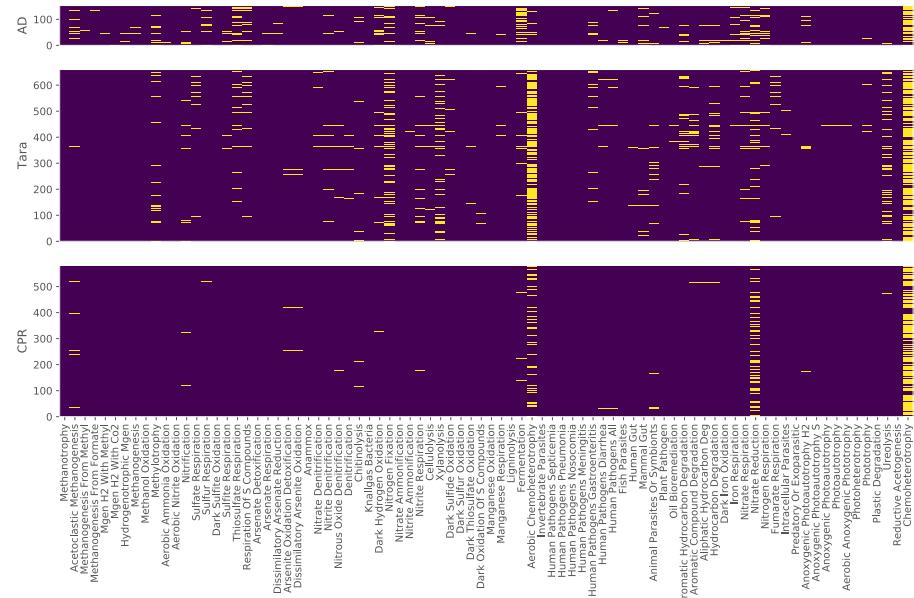
oceans project [5]), and MAGs from a groundwater aquifer assigned to be members of the so-called ‘candidate phyla radiation’ (CPR) [8]. The CPR is a set of bacterial lineages discovered from metagenomic studies consisting of a very large number of proposed novel phyla. These organisms have very small genomes, and are suggested to live in symbiosis with other organisms [42].

To perform the functional assignments, we used the  $\ell_1$ -regularized logistic regression classifier described above, with a random train-test splitting and the regularization parameter  $C = 0.05$ , trained using KEGG orthologs on the full NCBI genomes. Figure 5 shows a heatmap with presence or absence of the different traits for the MAGs assembled from anaerobic digesters and from the global oceans (the Tara project).

There are noticeable differences between traits assigned to MAGs from AD vs. ocean data, as shown in Figure S3 (and also Figure S3). Some of these differences reflect obvious differences in the two environments (e.g. aerobic traits missing from the AD data), while others point to more subtle differences that need further exploration.

These differences in trait might well be expected between these environments. For example, fermentation is very important in the AD process, and aerobic chemoheterotrophy obviously is not as the environment is aerobic. This indicates that the method is capable of producing useful information about MAGs. The results for the CPR MAGs indicate that these organisms possess significantly fewer traits than those from the other two environments, as would be expected from their very small genome sizes. A few traits do however have significant incidence in this group. Apart from ‘chemoheterotrophy’ and ‘aerobic chemoheterotrophy’, which are very broad categories encompassing a large proportion of all organisms, a few traits associated with nitrogen metabolism, especially nitrate reduction, are noticeably present in the group. That the CPR are involved in nitrate reduction was recently proposed in Danczak et al. [42].

Some of the functional assignments seem strange, for example organisms being classified as acetoclastic or hydrogenotrophic methanogens but not as methanogens, including a significant proportion of CPR organisms (about 3%), which are bacteria, being classified as acetoclastic methanogens. Looking at the gene orthologs present in these organisms and their taxonomic assignments sheds some light on what is going on here, see Table 3. For example, some of the acetoclastic methanogens which are misclassified as not being methanogens are missing the *mcrA* ortholog K00399, presumably because the MAGs are incomplete and this gene has been missed. Another



**Fig 5. Heatmap of presence/absence of fucntions in MAGs.** Results of running the set of logistic regression classifiers trained on NCBI genomes on MAGs assembled from three environments: laboratory anaerobic digesters, the ocean and ‘candidate phyla radiation’ (CPR) organisms from an aquifer system.

example is an organism classified as being a hydrogenotrophic methanogen but not a methanogen. This MAG appears to be from the bacterium *Caldisericum exile*, which is not a methanogen and does not possess *mcrA* (it is an anaerobic, thermophilic bacterium which respires by thiosulfate reduction ??). However, it does possess genes for subunits of the energy-converting hydrogenase A, noted earlier (Table 2) to be indicative of hydrogenotrophic methanogenesis. Therefore, these discrepancies may be the result either of incomplete MAGs, or of combinations of genes which are rare or unseen in the training set.

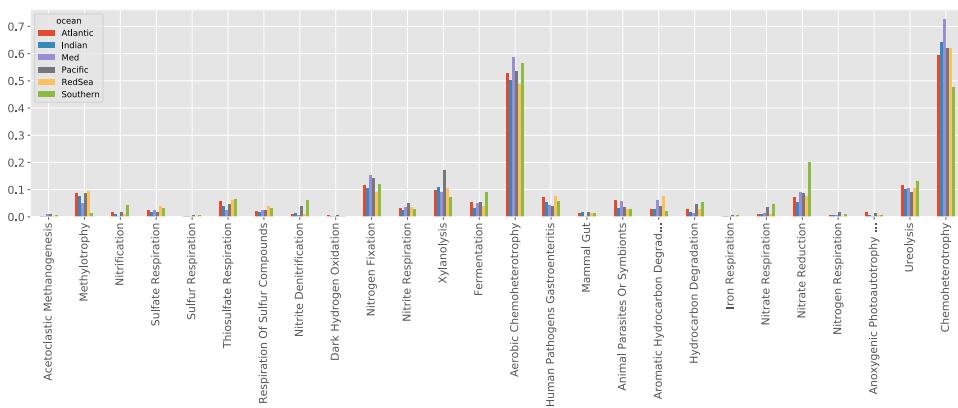
For the Tara dataset, metadata for different samples was available. Combining the coverages of the different MAGs across these samples with the functional assignments of the MAGs, we can calculate the proportion of total coverage which comes from microbes with a given trait. Figure 6 shows the mean of this metric for all the traits for samples from different oceans. The differences do not seem to be very significant between oceans, with a few exceptions, notably an abundance of nitrate reduction in the Southern Ocean. Samples metadata included depth, temperature and salinity, among other measurements. Figure 7 shows the correlation between temperature and fermenter abundance, showing a rather large negative correlation (Pearson correlation  $r = -0.43, p < 0.001$ ).

## Discussion

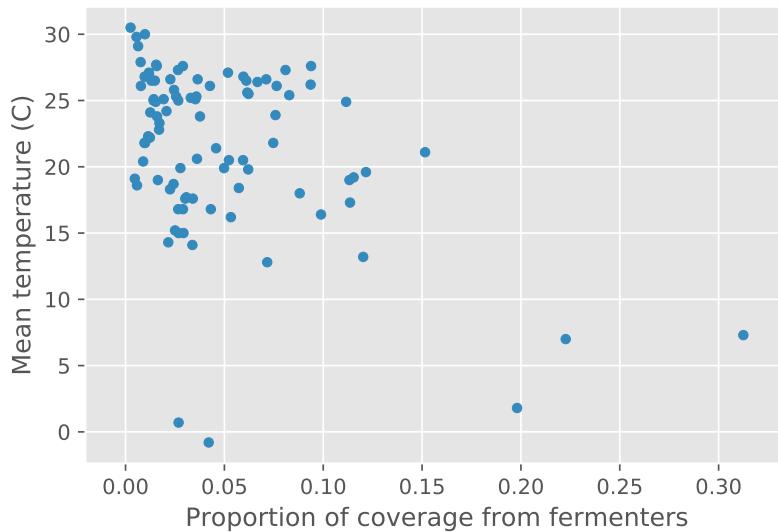
We have demonstrated a method for inferring phenotypes from genotypes, in the form of gene ortholog copy numbers, using machine learning on an existing phenotype database combined with NCBI genomes. While the accuracy of the predictions vary significantly over different phenotypes, a significant proportion of the traits we tested achieved very good classification accuracy, with AUROC scores greater than 90%. Of the machine learning algorithms we used, we found that  $\ell_1$ -regularized logistic regression gave the best combination of accuracy, computational efficiency and interpretability. The results

| Species                           | mcrA | echA | methanogenesis | hydrogenotrophic | acetoclastic |
|-----------------------------------|------|------|----------------|------------------|--------------|
| Caldisericum exile                | 0    | 1    | 0              | 1                | 0            |
| Methanomassiliicoccus luminyensis | 1    | 0    | 1              | 1                | 0            |
| Methanosaeta concilii             | 2    | 0    | 1              | 1                | 1            |
| Methanosaeta concilii             | 0    | 0    | 1              | 0                | 1            |
| Methanosaeta harundinacea         | 1    | 0    | 1              | 0                | 0            |
| Methanoregula formicata           | 1    | 0    | 1              | 0                | 0            |
| Methanoregula formicata           | 1    | 0    | 1              | 0                | 0            |
| Methanolinea tarda                | 1    | 0    | 1              | 0                | 0            |

**Table 3. Key genes and predicted traits for MAGs predicted to be methanogens.** Gene copy numbers for the *mcrA* methanogenesis gene and the energy-converting hydrgoenase A, along with functional predictions, for AD MAGs predicted to be methanogenic by our algorithm.



**Fig 6. Mean abundance of microbes performing traits by ocean.** Average of the proportion of total coverage associated with microbes possessing a trait over samples from each of the oceans.



**Fig 7. Association of temperature and fermenter abundance.** Scatterplot of proportion of total coverage associated with microbes performing fermentation versus mean temperature of a sample. Pearson  $r = -0.43, p < 0.001$ .

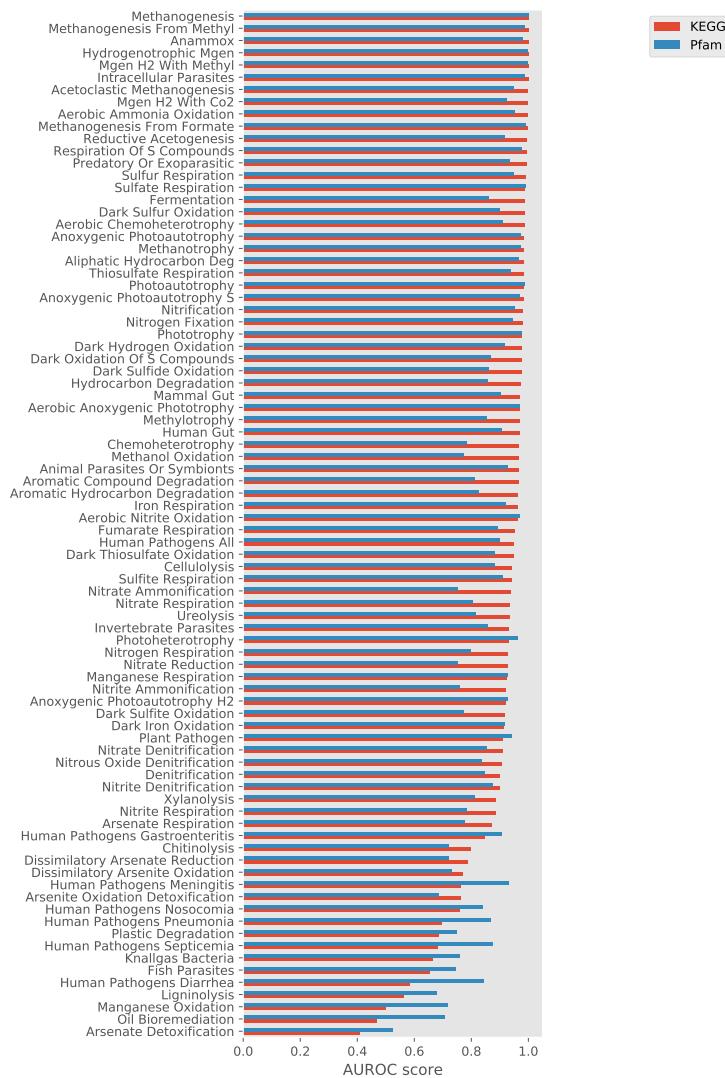
did not depend very strongly on whether KEGG orthologs or Pfam domains were used to characterise genes in the genomes, although the KEGG scheme performed slightly better on average over the traits we considered here. The logistic regression models we generate can be inspected, and the genes most associated with a given phenotype in the model enumerated, which allowed for validation of the models by comparison to what is known about the trait of these genes, and potentially for this method to discover new orthologous groups associated with phenotypes. For example, we found here that the presence of subunits of the energy-converting hydrogenase A are more predictive of an organism's performing hydrogenotrophic methanogenesis than the genes directly involved in the process as described in the KEGG functional module for it. For nitrate reduction, in addition to expected genes such as nitrate reductase, there are multiple KEGG orthologs listed as 'uncharacterised protein' which are highly predictive of the function.

To check the robustness of the models we generated, we tried training the models on one section of the microbial taxonomic tree, the Proteobacteria, and testing its accuracy on organisms from the rest of the tree, which it had not encountered at all in training. This did significantly reduce model accuracy for many phenotypes. This is to be expected, as the training and test sets in this case are so different. However, some of the traits still achieved good accuracy. This would suggest that the logistic regression model is identifying genes functionally involved with the phenotype in the training stage, such that their presence even in distantly related organisms is indicative of the presence of the phenotype. Phenotypes with good accuracy in this schema tended to produce models involving only a few genes (i.e. only a few genes had nonzero weights in the logistic regression model), less than 100, supporting the idea that they the models are picking out genes directly involved with the phenotype.

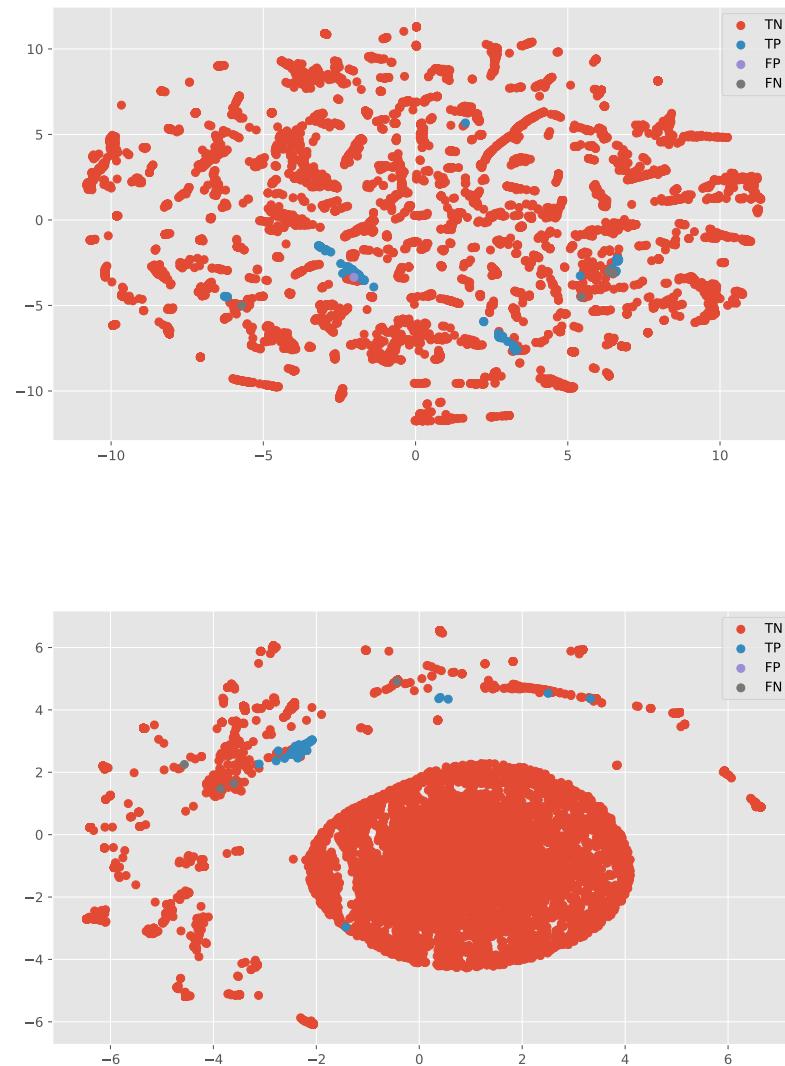
A major use we envisage for this method is the prediction of phenotypes of novel, uncultivated organisms discovered through metagenomic studies. To demonstrate this, we ran the trained models on collections of metagenomically assembled genomes (MAGs) from three environments: laboratory anaerobic digesters (ADs), the oceans (from the Tara oceans sequencing project), and 'candidate phyla radiation' MAGs from an aquifer system. There were significant differences in the predictions in the different environments, and these made sense in terms of what is known about the environments. For example, AD MAGs had high numbers of fermenters and low numbers assigned to 'aerobic chemoheterotrophy'. CPR MAGs had relatively few traits assigned overall, which makes sense given their very small genome size, but had a relatively large number of nitrate reducers [42]. Therefore, these trained models can help us to gain insights into the functional capabilities both of microbiomes as a whole and the individual species making up those communities, even when the species in these samples are novel and uncultivated.

## Supporting information

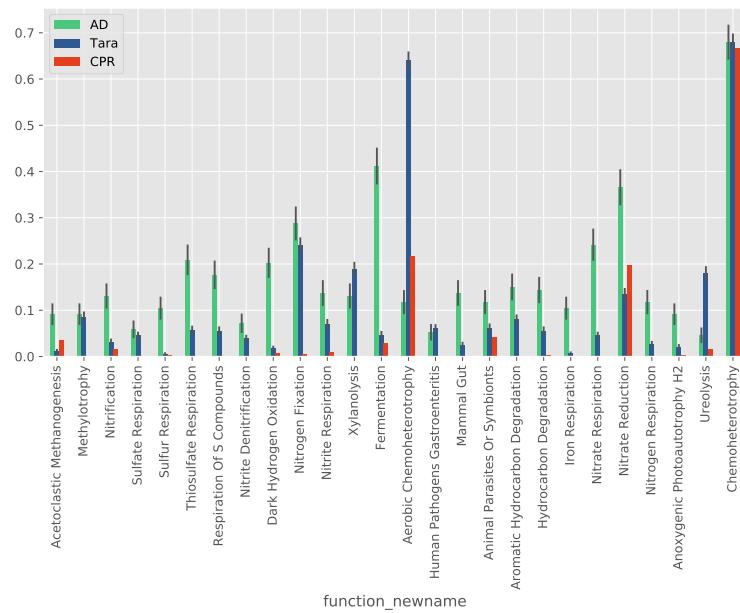
### Supplementary figures



**Fig S1. Performance of classification algorithms using different ortholog schemes.** The AUROC score on each classification task (each trait) is shown for algorithms trained on two representations of genomes in terms of orthologous groups of genes, the KEGG orthology (KO) and Pfam protein families.



**Fig S2. Stochastic neighbour embedding of all species in KEGG ortholog space** Ordination in two dimensions of all the species in the training dataset. By coloring by classification group (true positive, true negative, false positive, false negative) for a particular trait, here sulfate respiration, we can graphically visualise the behaviour of the classifier. Top: embedding performed over all KEGG orthologs. Bottom: embedding performed only over KOs relevant to the trait according to the classifier.



**Fig S3. Overall comparison of AD, Tara and CPR MAGs.** Proportions of MAGS from the environments having a trait, for some of the most common traits.

## References

1. Green JL, Bohannan BJM, Whitaker RJ. Microbial biogeography: from taxonomy to traits. *Science (New York, NY)*. 2008;320(5879):1039–43. doi:10.1126/science.1153475.
2. Martiny J BH, Jones SE, Lennon JT, Martiny AC. Microbiomes in light of traits: A phylogenetic perspective. *Science*. 2015;350(6261).
3. Huttenhower C, Gevers D, Knight R, Abubucker S, Badger JH, Chinwalla AT, et al. Structure, function and diversity of the healthy human microbiome. *Nature*. 2012;486(7402):207–214. doi:10.1038/nature11234.
4. Gilbert JA, Jansson JK, Knight R. The Earth Microbiome project: successes and aspirations. *BMC Biology*. 2014;12(1):69. doi:10.1186/s12915-014-0069-1.
5. Zhang H, Ning K. The Tara Oceans Project: New Opportunities and Greater Challenges Ahead. *Genomics, proteomics & bioinformatics*. 2015;13(5):275–7. doi:10.1016/j.gpb.2015.08.003.
6. Alneberg J, Bjarnason BS, de Brujin I, Schirmer M, Quick J, Ijaz UZ, et al. Binning metagenomic contigs by coverage and composition. *Nature Methods*. 2014;11(11):1144–1146. doi:10.1038/nmeth.3103.
7. Eren AM, Esen ÖC, Quince C, Vineis JH, Morrison HG, Sogin ML, et al. Anvi'o: an advanced analysis and visualization platform for ‘omics data. *PeerJ*. 2015;3:e1319. doi:10.7717/peerj.1319.
8. Anantharaman K, Brown CT, Hug LA, Sharon I, Castelle CJ, Probst AJ, et al. Thousands of microbial genomes shed light on interconnected biogeochemical processes in an aquifer system. *Nature Communications*. 2016;7:13219. doi:10.1038/ncomms13219.
9. Widder S, Allen RJ, Pfeiffer T, Curtis TP, Wiuf C, Sloan WT, et al. Challenges in microbial ecology: building predictive understanding of community function and dynamics. *The ISME Journal*. 2016;10:2557 EP –.
10. Louca S, Hawley AK, Katsev S, Torres-Beltran M, Bhatia MP, Kheirandish S, et al. Integrating biogeochemistry with multiomic sequence information in a model oxygen minimum zone. *Proceedings of the National Academy of Sciences*. 2016;113(40):E5925–E5933. doi:10.1073/pnas.1602897113.
11. Philippot L, Bru D, Saby NPA, Čuhel J, Arrouays D, Šimek M, et al. Spatial patterns of bacterial taxa in nature reflect ecological traits of deep branches of the 16S rRNA bacterial tree. *Environmental Microbiology*. 2009;11(12):3096–3104. doi:10.1111/j.1462-2920.2009.02014.x.
12. Louca S, Parfrey LW, Doebeli M. Decoupling function and taxonomy in the global ocean microbiome. *Science (New York, NY)*. 2016;353(6305):1272–7. doi:10.1126/science.aaf4507.
13. Louca S, Doebeli M. Reaction-centric modeling of microbial ecosystems. *Ecological Modelling*. 2016;335:74–86. doi:10.1016/J.ECOLMODEL.2016.05.011.
14. Philippot L, Andersson SGE, Battin TJ, Prosser JI, Schimel JP, Whitman WB, et al. The ecological coherence of high bacterial taxonomic ranks. *Nature Reviews Microbiology*. 2010;8(7):523–529. doi:10.1038/nrmicro2367.

15. SABARLY V, BOUDET O, GLODT J, CLERMONT O, SKURNIK D, DIANCOURT L, et al. The decoupling between genetic structure and metabolic phenotypes in *Escherichia coli* leads to continuous phenotypic diversity. *Journal of Evolutionary Biology*. 2011;24(7):1559–1571. doi:10.1111/j.1420-9101.2011.02287.x.
16. Koeppel AF, Wu M. Lineage-dependent ecological coherence in bacteria. *FEMS Microbiology Ecology*. 2012;81(3):574–582. doi:10.1111/j.1574-6941.2012.01387.x.
17. Martiny AC, Treseder K, Pusch G. Phylogenetic conservatism of functional traits in microorganisms. *The ISME Journal*. 2013;7(4):830–838. doi:10.1038/ismej.2012.160.
18. Narayanasamy S, Jarosz Y, Muller EEL, Heintz-Buschart A, Herold M, Kaysen A, et al. IMP: a pipeline for reproducible reference-independent integrated metagenomic and metatranscriptomic analyses. *Genome biology*. 2016;17(1):260. doi:10.1186/s13059-016-1116-8.
19. Kanehisa M, Furumichi M, Tanabe M, Sato Y, Morishima K. KEGG: new perspectives on genomes, pathways, diseases and drugs. *Nucleic Acids Research*. 2017;45(D1):D353–D361. doi:10.1093/nar/gkw1092.
20. Finn RD, Coggill P, Eberhardt RY, Eddy SR, Mistry J, Mitchell AL, et al. The Pfam protein families database: towards a more sustainable future. *Nucleic Acids Research*. 2016;44(D1):D279–D285. doi:10.1093/nar/gkv1344.
21. Tatusov RL, Koonin EV, Lipman DJ. A genomic perspective on protein families. *Science (New York, NY)*. 1997;278(5338):631–7.
22. Louca S, Jacques SMS, Pires APF, Leal JS, Srivastava DS, Parfrey LW, et al. High taxonomic variability despite stable functional structure across microbial communities. *Nature ecology & evolution*. 2016;1(1):15. doi:10.1038/s41559-016-0015.
23. Louca S, Jacques SMS, Pires APF, Leal JS, González AL, Doebeli M, et al. Functional structure of the bromeliad tank microbiome is strongly shaped by local geochemical conditions. *Environmental Microbiology*. 2017;19(8):3132–3151. doi:10.1111/1462-2920.13788.
24. Łukaszewicz M, Jabłoński S, Rodowicz P. Methanogenic archaea database containing physiological and biochemical characteristics. *International Journal of Systematic and Evolutionary Microbiology*. 2015;65(4):1360–1368. doi:10.1099/ij.s.0.000065.
25. Whitman WB, Bergey's Manual Trust, Wiley Online Library (Online service). Bergey's manual of systematics of archaea and bacteria;.
26. Weimann A, Mooren K, Frank J, Pope PB, Bremges A, McHardy AC. From genomes to phenotypes: Traitar, the microbial trait analyzer; 2016. Available from: <http://biorxiv.org/lookup/doi/10.1101/043315>.
27. Berger SA. GIDEON: a comprehensive Web-based resource for geographic medicine. *International Journal of Health Geographics*. 2005;4(1):10. doi:10.1186/1476-072X-4-10.
28. Hyatt D, Chen GL, LoCascio PF, Land ML, Larimer FW, Hauser LJ. Prodigal: prokaryotic gene recognition and translation initiation site identification. *BMC Bioinformatics*. 2010;11(1):119. doi:10.1186/1471-2105-11-119.

29. Buchfink B, Xie C, Huson DH. Fast and sensitive protein alignment using DIAMOND. *Nature Methods*. 2014;12(1):59–60. doi:10.1038/nmeth.3176.
30. Eddy SR. Accelerated Profile HMM Searches. *PLoS Computational Biology*. 2011;7(10):e1002195. doi:10.1371/journal.pcbi.1002195.
31. Hastie T, Tibshirani R, Friedman J. Linear Methods for Classification. In: *The Elements of Statistical Learning*; 2009. p. 101–137.
32. Freedman D. Statistical models : theory and practice. Cambridge University Press; 2009. Available from: <http://www.cambridge.org/gb/academic/subjects/statistics-probability/statistical-theory-and-methods/statistical-models-theory-and-practice-2nd-edition?format=HB{&}isbn=9780521112437{#}IqXivBILqUALXPKZ.97>.
33. Lee SI, Lee H, Abbeel P, Ng AN. Efficient L1 Regularized Logistic Regression. In: *The Twenty-First National Conference on Artificial Intelligence and the Eighteenth Innovative Applications of Artificial Intelligence Conference*; 2006.
34. Hastie T, Tibshirani R, Friedman J. Random Forests. In: *The Elements of Statistical Learning*; 2009. p. 1–18.
35. Fawcett T, Tom. An introduction to ROC analysis. *Pattern Recognition Letters*. 2006;27(8):861–874. doi:10.1016/j.patrec.2005.10.010.
36. Flach P, Hernández-Orallo J, Ferri C. A coherent interpretation of auc as a measure of aggregated classification performance. *Proceedings of the 28th International Conference on Machine Learning*, Bellevue, WA, USA. 2011; p. 657–664. doi:10.1145/347090.347126.
37. Delmont TO, Quince C, Shaiber A, Esen OC, Lee STM, Lucker S, et al. Nitrogen-Fixing Populations Of Planctomycetes And Proteobacteria Are Abundant In The Surface Ocean. *bioRxiv*. 2017; p. 129791. doi:10.1101/129791.
38. Sunagawa S, Coelho LP, Chaffron S, Kultima JR, Labadie K, Salazar G, et al. Ocean plankton. Structure and function of the global ocean microbiome. *Science* (New York, NY). 2015;348(6237):1261359. doi:10.1126/science.1261359.
39. Brown CT, Hug LA, Thomas BC, Sharon I, Castelle CJ, Singh A, et al. Unusual biology across a group comprising more than 15% of domain Bacteria. *Nature*. 2015;523(7559):208–211. doi:10.1038/nature14486.
40. U D. The unique biochemistry of methanogenesis. *Prog Nucleic Acid Res Mol Biol*. 2002;71:223–83.
41. Chan L, Morgan-Kiss RM, Hanson TE. Functional Analysis of Three Sulfide:Quinone Oxidoreductase Homologs in *Chlorobaculum tepidum*. *J Bacteriol*. 2009;10:1026–1034.
42. Danczak RE, Johnston MD, Kenah C, Slattery M, Wrighton KC, Wilkins MJ. Members of the Candidate Phyla Radiation are functionally differentiated by carbon- and nitrogen-cycling capabilities. *Microbiome*. 2017;5:112.